



# YOLOPilot

## **Self driving through YOLOv8 and Transfer learning**

---

Presented by: Mohammad  
Reza Azadi Tinat

# Project Definition

## Objective

To develop an object detection system for autonomous driving using YOLOv8, capable of identifying key objects such as:

- Motorbikes
- Vehicles
- Pedestrians
- Traffic Lights

## Core Idea

Train and evaluate a YOLO-based deep learning model on two different real-world datasets, applying both training from scratch and transfer learning to compare performance and efficiency.

## Goals

- Accurately detect traffic-relevant objects in road scenes
- Compare different training strategies (e.g., full training vs. transfer learning)
- Build an interactive **Gradio-based app** for testing detection

# Application steps



## 1. Dataset Preparation

- Two datasets were organized into the YOLOv8 format: train/images, train/labels, etc.
- Label files were cleaned and converted for consistency with a unified class schema.

## 2. Initial Training on Dataset 1

- A YOLOv8 base model (yolov8s.pt) was trained on the first dataset.
- Performance was evaluated using F1 scores and confusion matrices.

## 3. Error Analysis

- Detected mismatches between predictions and ground truth.
- Visualized incorrect detections to identify model limitations.

## 4. Transfer Learning on Dataset 2

- Best weights from Dataset 1 were used to fine-tune on the second dataset.
- Experiments included resized and padded image inputs.

## 5. Evaluation and Visualization

- Metrics, confusion matrices, and sample results were collected.
- Validated model generalization and identified areas for improvement.

# Project Structure

```
YOLOPilot/
├── first_dataset/
│   ├── train
│   │   ├── images
│   │   ├── lables_old
│   │   └── lables
│   ├── valid
│   │   ├── images
│   │   ├── lables_old
│   │   └── lables
│   └── test
│       ├── images
│       ├── lables_old
│       └── lables
├── runs → Results after training
├── second_dataset/
│   ├── train(Resized), train_orignal
│   │   ├── images
│   │   └── lables
│   ├── val(Resized), val_orignal
│   │   ├── images
│   │   └── lables
│   ├── test(Resized), test_orignal
│   │   ├── images
│   │   └── lables
│   ├── labels_train.csv, labels_val.csv
│   ├── new_labels
│   ├── images_resized
│   ├── resized_labels
│   └── runs → Results after training
├── src/
│   ├── label_modifier.py
│   ├── second_dataset_resizer.py
│   └── second_dataset_splitter.py
├── docs/
│   ├── images/
│   │   ├── logo.png
│   │   └── app_screenshot.png
├── PostProcessing_box.py
├── PreProcessing.py
├── train_total.py
├── Dockerfile
├── requirements.txt
├── yolopilot_app.py
├── Dockerfile
└── README.md
```

- The **first\_dataset/** and **second\_dataset/** folders store the two main datasets, each with separate folders for images, labels\_old, and modified labels.
- Script **PreProcessing.py** is used during data preprocessing.
- The root directory has all the essential scripts for training (**train\_total.py**), validation and post-processing (**PostProcessing\_box.py**), and deployment files like the Dockerfile, requirements.txt, and **yolopilot\_app.py** for running the Gradio interface.
- This structure keeps the project clean, scalable, and easy to reproduce.

# Dataset 1 – CARLA Object Detection

**Source:** [CARLA-based Self-Driving Dataset](#)

**Environment:** Simulated driving scenes from the CARLA simulator

**Classes:**

- bike, motobike, person,
- traffic\_light\_\*, traffic\_sign\_\*, vehicle

**Format:** YOLOv5-style (images + labels in .txt)

**Label Issues:**

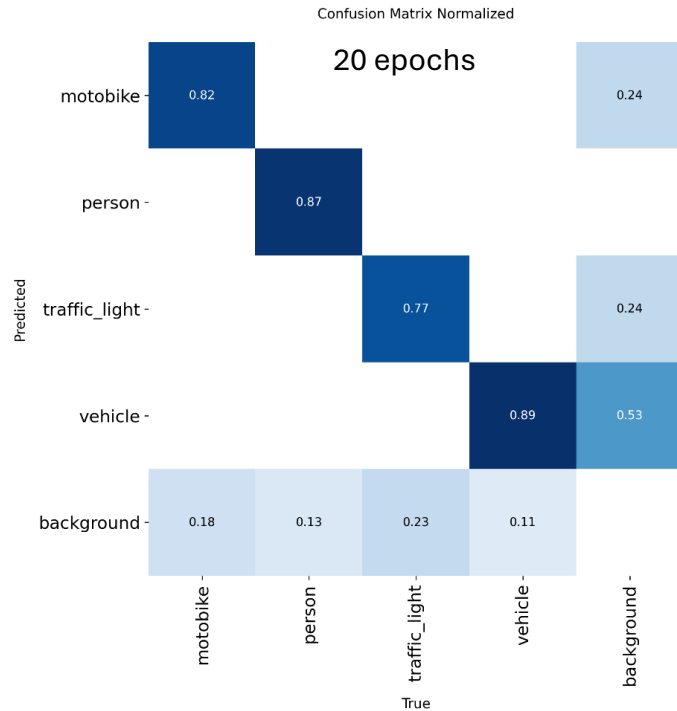
- Redundant categories (e.g., multiple traffic light classes)
- Needed mapping and simplification before training

**Training:**

- Conducted 3 separate training runs with different epochs (20,50,100)
- One additional training via *continued training* from a previous run
- Model: yolov8s.pt

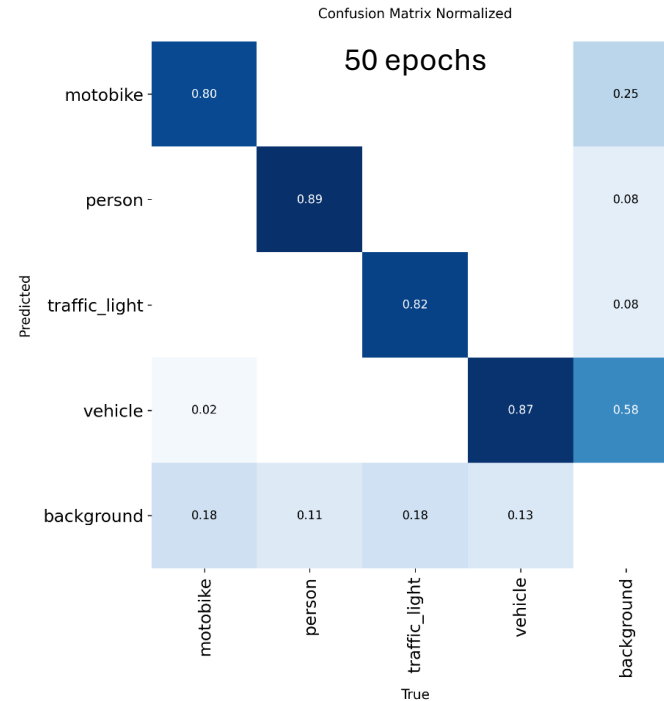


# Confusion Matrix Comparison – Different Epochs (Dataset 1)



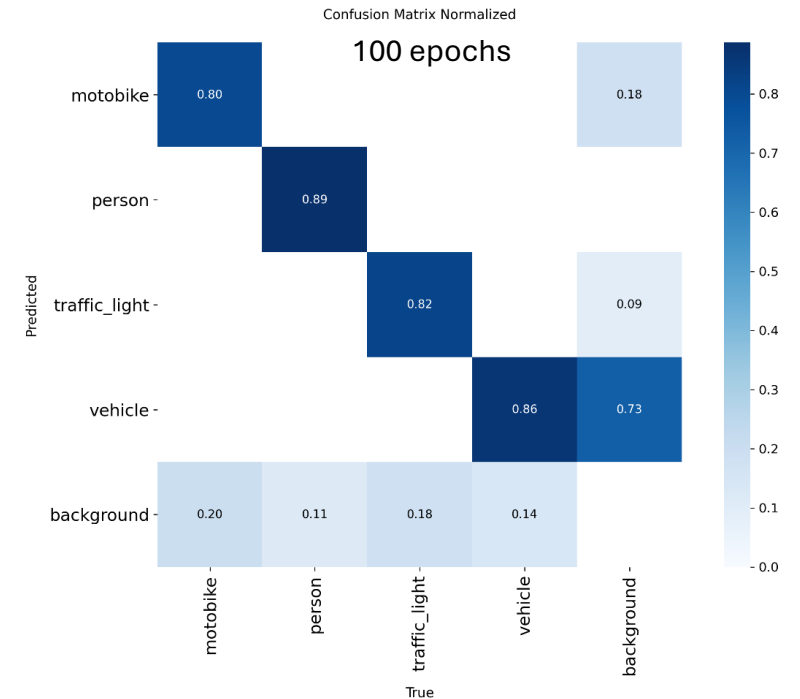
## Epoch 20:

- High confusion between traffic\_light and background (e.g., 23% misclassified)
- Early-stage learning with weak generalization



## Epoch 50:



- Strong improvement in class separation
- High accuracy for pedestrian and vehicle



## Epoch 100:

- Class-wise accuracy remains stable
- Gains are marginal compared to 50 epochs

# Continued Training (Warm Start)



- Minor improvement in traffic\_light (0.83 vs. 0.82)
- Drop in vehicle precision (0.81 vs. 0.86)
- Increased background confusion
- No significant gain over 50-epoch standalone model

## Conclusion:

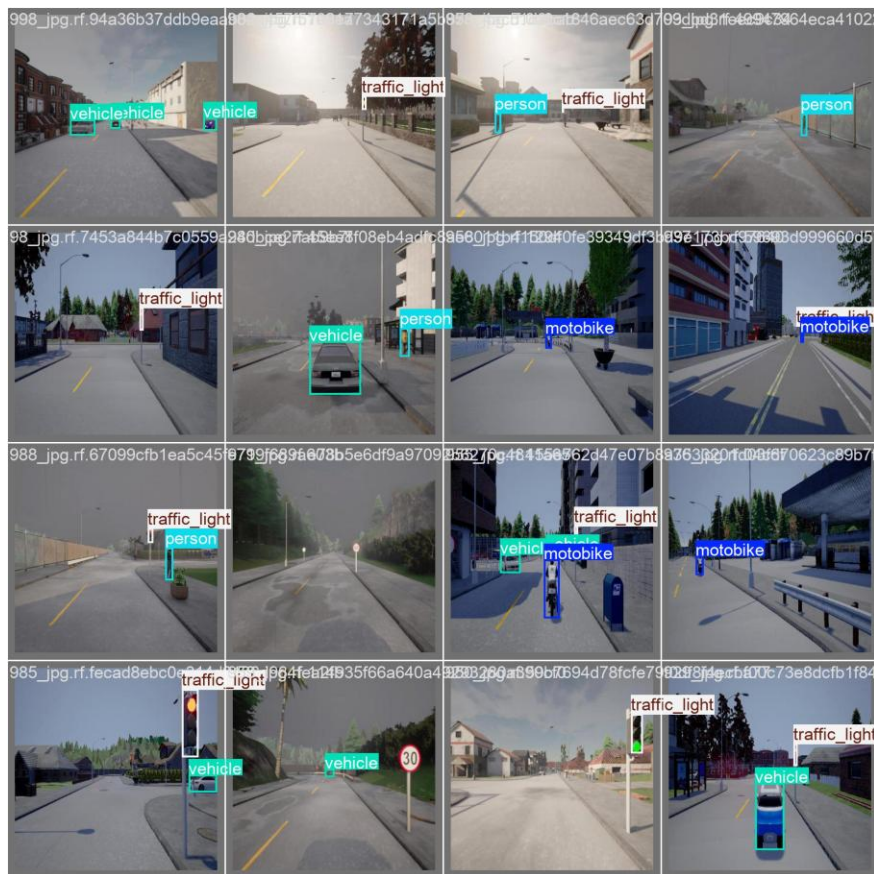
Training plateau reached after 50 epochs

→ Continued training provided diminishing returns

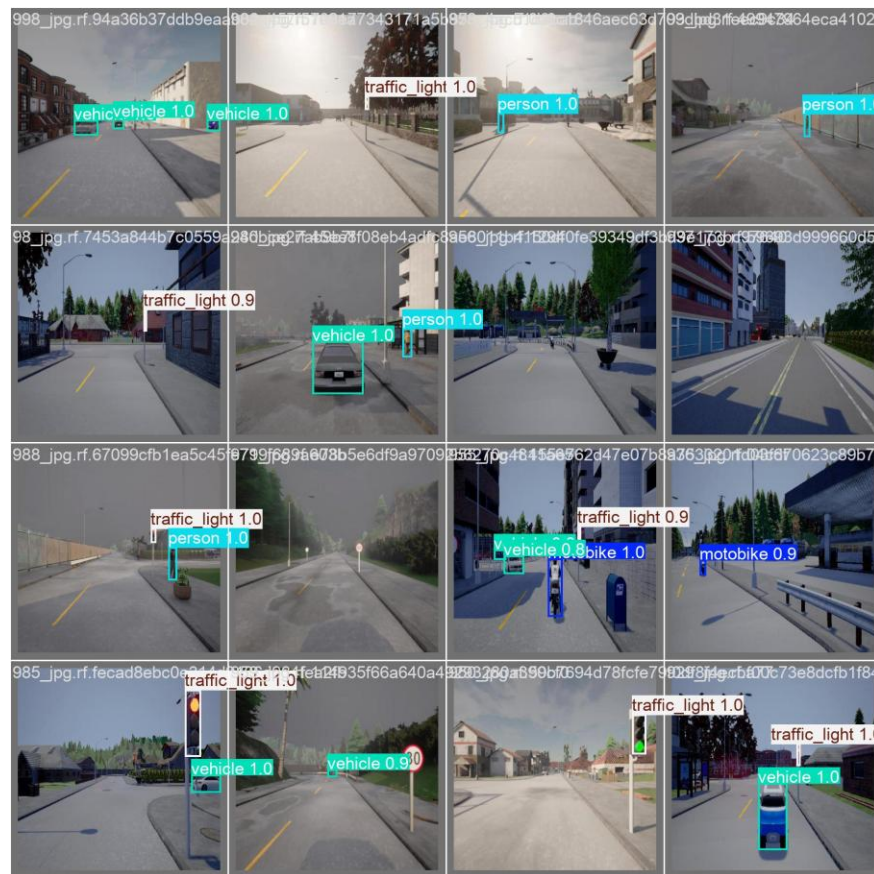
→ 50-epoch model selected for transfer learning on Dataset 2



# Sample Visual Results – Dataset 1 (50 Epochs)



Actual labels from validation images



Predicted labels from validation images

## The model demonstrates:



- Strong performance on vehicle, person, and motobike detection.
- Tight bounding boxes and high confidence scores for clearly visible objects.

## Challenges observed:

- Occasional missed detections of traffic lights, especially when small, occluded, or poorly lit.
- Some misclassifications (e.g., traffic light predicted as person, or missed motobike in crowded scenes).



# Evaluation on Second Dataset (Before Transfer Learning)



- Dataset 2 features real-world traffic scenes from video frames — introducing challenges like:
  - Varying lighting conditions
  - Occlusions and clutter
  - Diverse object sizes and angles

## Initial Results (Without Transfer Learning)

- The original model fails to generalize to this new domain.
- Most detections are classified as background, with key classes like *traffic\_light* and *motobike* frequently missed.
- The domain shift reveals that training only on synthetic or limited environments is insufficient for robust real-world deployment.

# Second Dataset Setup

To adapt the YOLOv8 model to the second (real-world) dataset, two preprocessing strategies were tested.

Both setups used best.pt weights from 50-epoch training on the first dataset as the starting point.

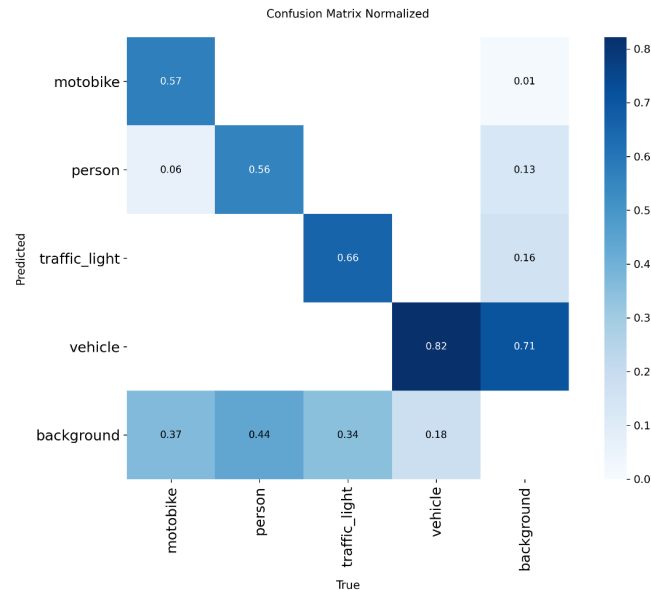
## **Approach 1: Explicit Resizing (416×416)**

- All images resized from **480×300** → **416×416**

## **Approach 2: Auto Resize & Padding (YOLOv8 Default - 640×640)**

- Original images passed without resizing
- YOLOv8 internally pads and resizes to 640×640

# Approach 1: Explicit Resizing




## Confusion Matrix Insights (Figure 13)

- Vehicle detection is highly accurate (0.82)
- Motobike and Traffic Light classes achieve moderate performance (0.57, 0.66)
- Person class remains difficult (0.56)
- Clear improvement from background-dominated misclassifications seen prior to transfer learning
- ✓ Confirms manual resizing + transfer learning is effective in adapting to real-world scenes

- Red boxes = Model predictions; Green boxes = Ground truth
- Objects such as vehicles, traffic lights, and pedestrians are successfully detected
- Significant improvement over pre-transfer model, which often missed all objects
- Minor issues:
  - Some duplicate or overlapping boxes
  - Occasional misalignment between prediction and actual object



# Approach 1: Explicit Resizing



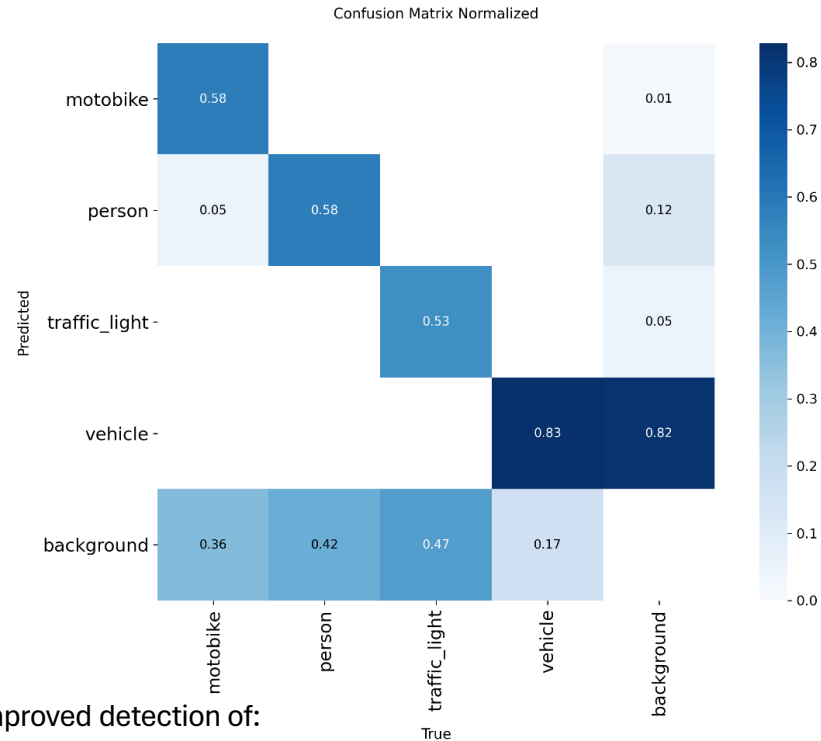
Actual labels from validation images



Predicted labels from validation images

- Accurate and consistent detection of multiple vehicles in real traffic scenes
- Tight bounding boxes with confidence scores above 0.9 for most vehicles
- Occasional successful pedestrian detections, proving class generalization

# Approach 2: Auto Resize & Padding



Improved detection of:

- **Persons** – thanks to better spatial preservation
- **Motobikes** and **vehicles** – aided by higher input resolution
- Slight drop in **traffic light** accuracy
- Padding may reduce visibility of small-scale or distant objects
- Background confusion reduced, indicating clearer distinction between object and non-object areas



## Conclusion:

- This method offers enhanced generalization for larger objects
- Performs slightly better overall than Approach 1 for complex, real-world scenes
- However, manual resizing may still be preferable for small object detection like traffic lights



# Approach 2: Auto Resize & Padding





Actual labels from validation images



Predicted labels from validation images

# Training on second data set without transfer learning



- Improved class separation overall
- Motobike and traffic\_light accuracy slightly better than with transfer learning
- Person detection more consistent
- Vehicle detection remains strong and comparable
- Confusion matrix shows cleaner boundaries between object classes and background

- Training time: 5h 30m (vs. 7h with transfer learning)
- Transfer learning required extra adaptation overhead



# Conclusion and Further Work

## Summary of Study

Evaluated multiple training strategies on a real-world object detection task using YOLOv8:

- Transfer Learning – Resized with label adjustment
- Transfer Learning – YOLO's Auto Resize & Padding
- Training from Scratch on Dataset 2

## Key Findings

- Transfer learning improves generalization over no pretraining
- Manual resizing + corrected labels performs better than auto-padding for small objects
- Training from scratch gives best accuracy and fastest training time.
  - Especially for motobike and traffic\_light classes

## Conclusion

- YOLOv8s can be effectively trained from scratch on real-world data. Transfer learning helps, but is not always necessary when dataset size and quality are sufficient.

## Future Directions

- Experiment with larger models: YOLOv8m, YOLOv8l



Thank you for  
your attention

# YOLOPilot