



YOLOPilot

Self driving through YOLOv8 and Transfer learning

Presented by: Mohammad
Reza Azadi Tinat

Project Definition

Objective

To develop an object detection system for autonomous driving using YOLOv8, capable of identifying key objects such as:

- Motorbikes
- Vehicles
- Pedestrians
- Traffic Lights

Core Idea

Train and evaluate a YOLO-based deep learning model on two different real-world datasets, applying both training from scratch and transfer learning to compare performance and efficiency.

Goals

- Accurately detect traffic-relevant objects in road scenes
- Compare different training strategies (e.g., full training vs. transfer learning)
- Build an interactive **Gradio-based app** for testing detection

Application steps



1. Dataset Preparation

- Two datasets were organized into the YOLOv8 format: train/images, train/labels, etc.
- Label files were cleaned and converted for consistency with a unified class schema.

2. Initial Training on Dataset 1

- A YOLOv8 base model (yolov8s.pt) was trained on the first dataset.
- Performance was evaluated using F1 scores and confusion matrices.

3. Error Analysis

- Detected mismatches between predictions and ground truth.
- Visualized incorrect detections to identify model limitations.

4. Transfer Learning on Dataset 2

- Best weights from Dataset 1 were used to fine-tune on the second dataset.
- Experiments included resized and padded image inputs.

5. Evaluation and Visualization

- Metrics, confusion matrices, and sample results were collected.
- Validated model generalization and identified areas for improvement.

Project Structure

- The **first_dataset/** and **second_dataset/** folders store the two main datasets, each with separate folders for images, labels_old, and modified labels.
- Scripts such as **label_modify.py**, **resize.py**, and **split.py** are used during data preprocessing.
- The root directory has all the essential scripts for training (**train_total.py**), validation (**valid_total.py**), post-processing (**PostProcessing_box.py**), and deployment files like the Dockerfile, requirements.txt, and **yolopilot_app.py** for running the Gradio interface.
- This structure keeps the project clean, scalable, and easy to reproduce.

```
YOLOPilot/
├── first_dataset/
│   ├── train
│   │   ├── images
│   │   ├── labels_old
│   │   └── labels
│   ├── valid
│   │   ├── images
│   │   ├── labels_old
│   │   └── labels
│   └── test
│       ├── images
│       ├── labels_old
│       └── labels
├── label_modify.py
├── runs → Results after training
├── second_dataset/
│   ├── train(Resized), train_orignal
│   │   ├── images
│   │   └── labels
│   ├── val(Resized), val_orignal
│   │   ├── images
│   │   └── labels
│   ├── test(Resized), test_orignal
│   │   ├── images
│   │   └── labels
│   ├── labels_train.csv, labels_val.csv
│   ├── new_labels
│   ├── images_resized
│   ├── resized_labels
│   ├── label_modify.py
│   ├── resize.py
│   ├── split.py
│   └── runs → Results after training
├── docs/
│   ├── images/
│   │   ├── logo.png
│   │   └── app_screenshot.png
├── PostProcessing_box.py
├── train_total.py
├── valid_total.py
├── Dockerfile
├── requirements.txt
├── yolopilot_app.py
├── Dockerfile
└── README.md
```

Dataset 1 – CARLA Object Detection

Source: [CARLA-based Self-Driving Dataset](#)

Environment: Simulated driving scenes from the CARLA simulator

Classes:

- bike, motobike, person,
- traffic_light_*, traffic_sign_*, vehicle

Format: YOLOv5-style (images + labels in .txt)

Label Issues:

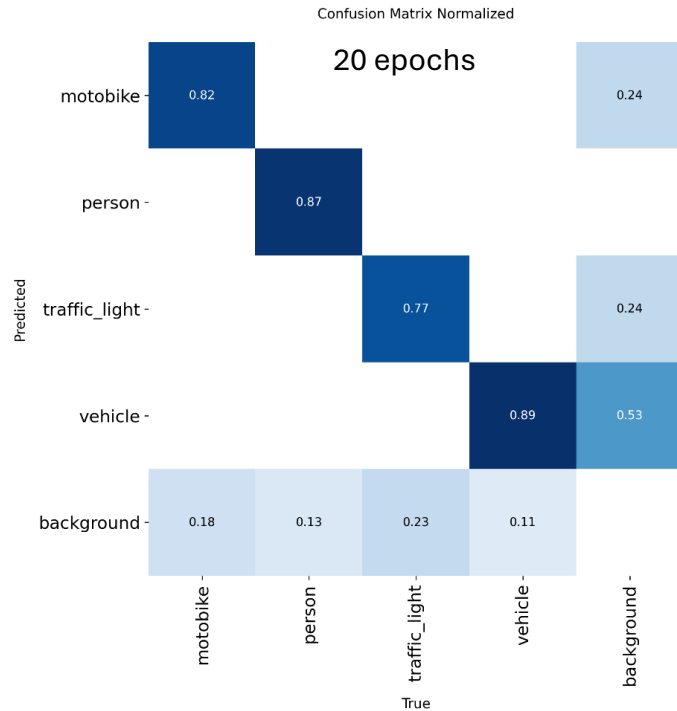
- Redundant categories (e.g., multiple traffic light classes)
- Needed mapping and simplification before training

Training:

- Conducted 3 separate training runs with different epochs (20,50,100)
- One additional training via *continued training* from a previous run
- Model: yolov8s.pt




Confusion Matrix Comparison – Different Epochs (Dataset 1)



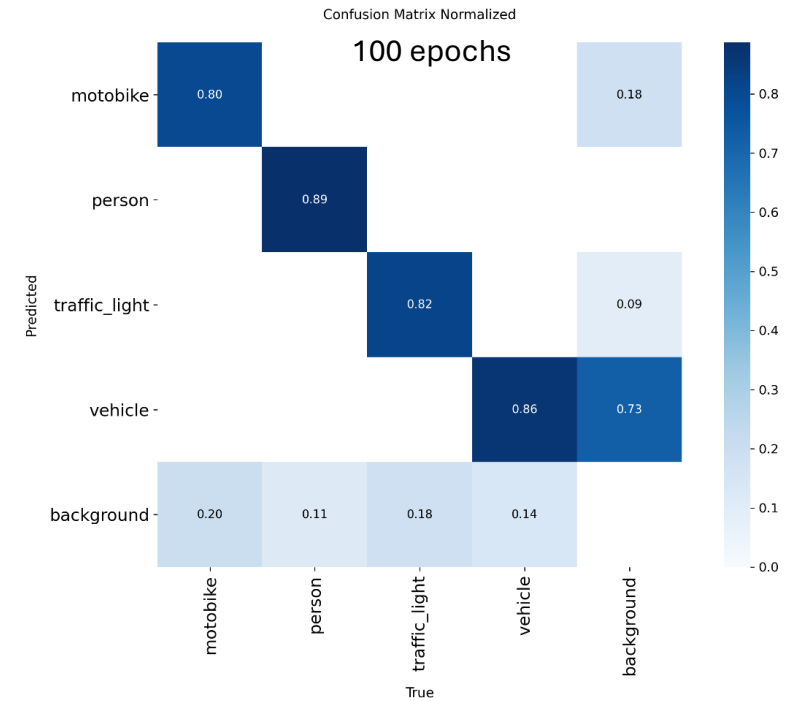
Epoch 20:

- High confusion between traffic_light and background (e.g., 23% misclassified)
- Early-stage learning with weak generalization



Epoch 50:



- Strong improvement in class separation
- High accuracy for pedestrian and vehicle



Epoch 100:

- Class-wise accuracy remains stable
- Gains are marginal compared to 50 epochs

Continued Training (Warm Start)



- Minor improvement in traffic_light (0.83 vs. 0.82)
- Drop in vehicle precision (0.81 vs. 0.86)
- Increased background confusion
- No significant gain over 50-epoch standalone model

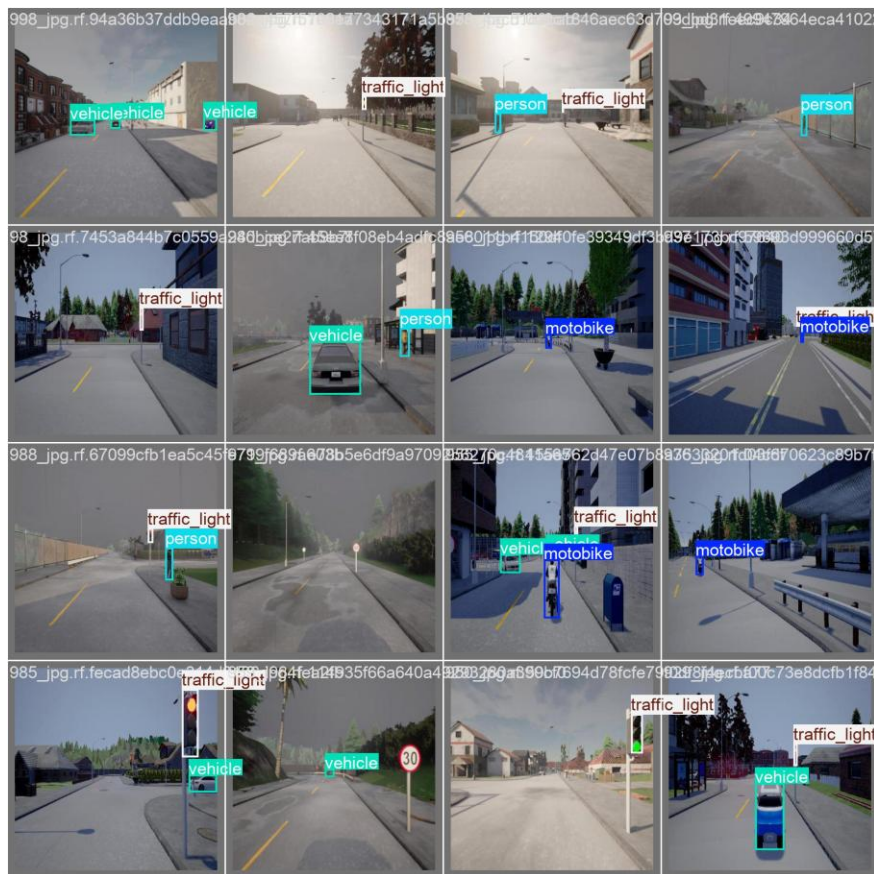
Conclusion:

Training plateau reached after 50 epochs

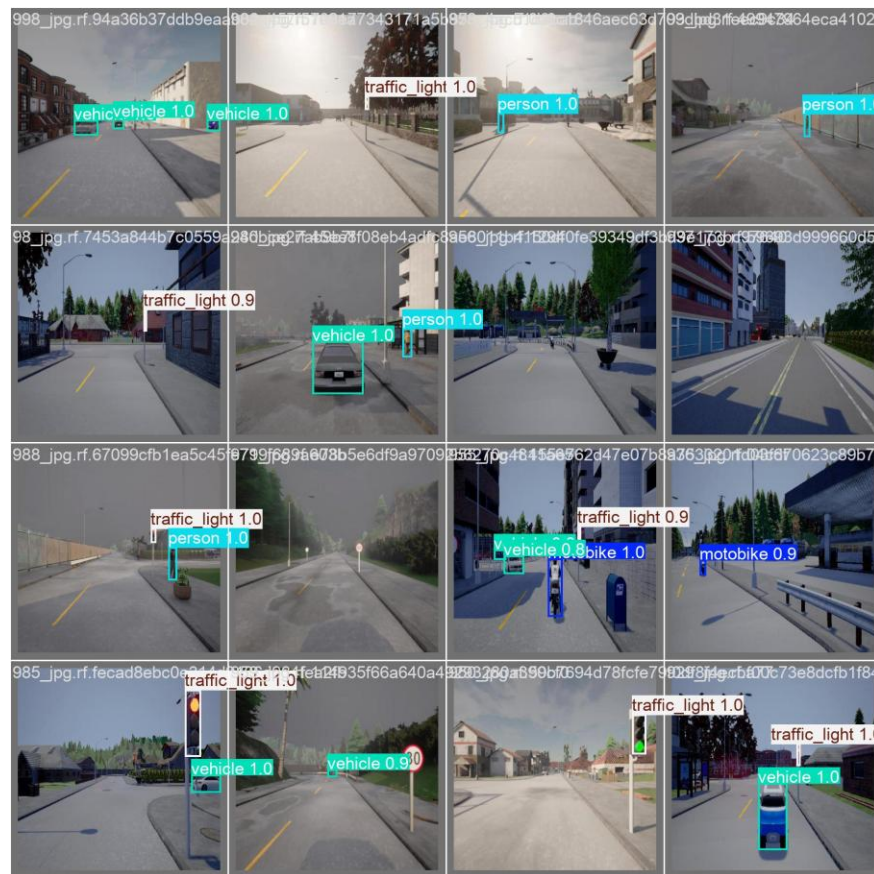
→ Continued training provided diminishing returns

→ 50-epoch model selected for transfer learning on Dataset 2

Sample Visual Results – Dataset 1 (50 Epochs)



Actual labels from validation images



Predicted labels from validation images



The model demonstrates:

- Strong performance on vehicle, person, and motobike detection.
- Tight bounding boxes and high confidence scores for clearly visible objects.

Challenges observed:

- Occasional missed detections of traffic lights, especially when small, occluded, or poorly lit.
- Some misclassifications (e.g., traffic light predicted as person, or missed motobike in crowded scenes).

Evaluation on Second Dataset (Before Transfer Learning)



- Dataset 2 features real-world traffic scenes from video frames — introducing challenges like:
 - Varying lighting conditions
 - Occlusions and clutter
 - Diverse object sizes and angles

Initial Results (Without Transfer Learning)

- The original model fails to generalize to this new domain.
- Most detections are classified as background, with key classes like *traffic_light* and *motobike* frequently missed.
- The domain shift reveals that training only on synthetic or limited environments is insufficient for robust real-world deployment.

Second Dataset Setup

To adapt the YOLOv8 model to the second (real-world) dataset, two preprocessing strategies were tested.

Both setups used best.pt weights from 50-epoch training on the first dataset as the starting point.

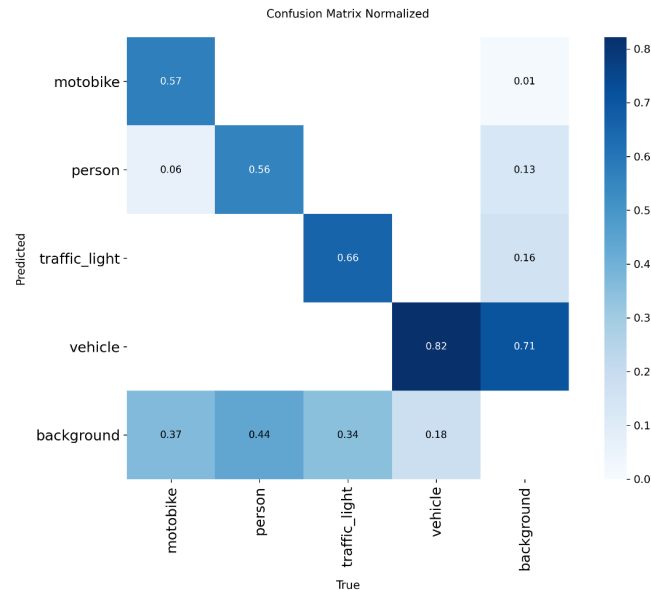
Approach 1: Explicit Resizing (416×416)

- All images resized from **480×300** → **416×416**

Approach 2: Auto Resize & Padding (YOLOv8 Default - 640×640)

- Original images passed without resizing
- YOLOv8 internally pads and resizes to 640×640

Approach 1: Explicit Resizing




Confusion Matrix Insights (Figure 13)

- Vehicle detection is highly accurate (0.82)
- Motobike and Traffic Light classes achieve moderate performance (0.57, 0.66)
- Person class remains difficult (0.56)
- Clear improvement from background-dominated misclassifications seen prior to transfer learning
- ✓ Confirms manual resizing + transfer learning is effective in adapting to real-world scenes

- Red boxes = Model predictions; Green boxes = Ground truth
- Objects such as vehicles, traffic lights, and pedestrians are successfully detected
- Significant improvement over pre-transfer model, which often missed all objects
- Minor issues:
 - Some duplicate or overlapping boxes
 - Occasional misalignment between prediction and actual object

Approach 1: Explicit Resizing



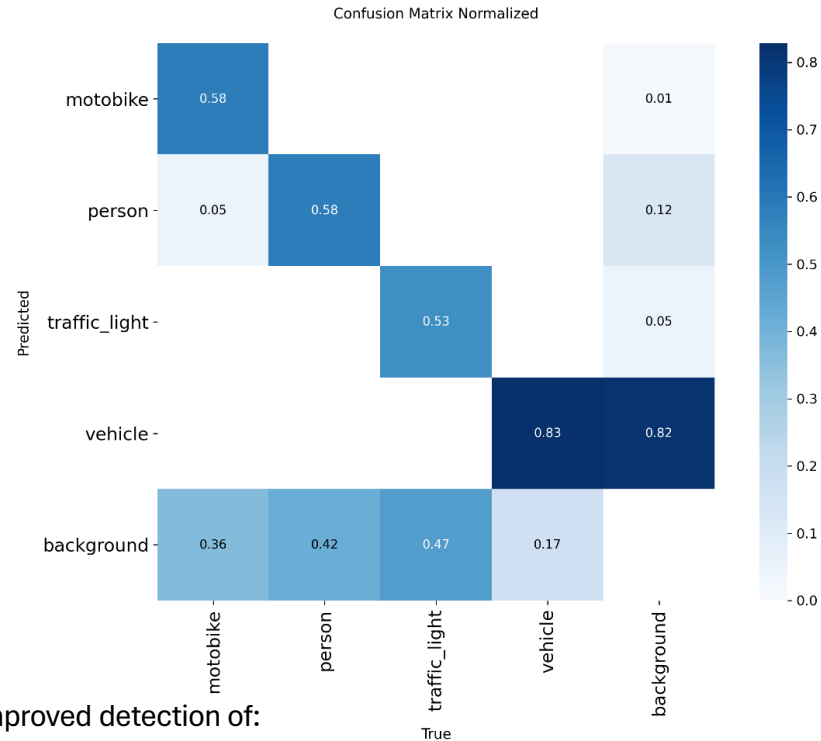
Actual labels from validation images



Predicted labels from validation images

- Accurate and consistent detection of multiple vehicles in real traffic scenes
- Tight bounding boxes with confidence scores above 0.9 for most vehicles
- Occasional successful pedestrian detections, proving class generalization

Approach 2: Auto Resize & Padding



Improved detection of:

- **Persons** – thanks to better spatial preservation
- **Motobikes** and **vehicles** – aided by higher input resolution
- Slight drop in **traffic light** accuracy
- Padding may reduce visibility of small-scale or distant objects
- Background confusion reduced, indicating clearer distinction between object and non-object areas




Conclusion:

- This method offers enhanced generalization for larger objects
- Performs slightly better overall than Approach 1 for complex, real-world scenes
- However, manual resizing may still be preferable for small object detection like traffic lights

Approach 2: Auto Resize & Padding





Actual labels from validation images



Predicted labels from validation images

Training on second data set without transfer learning



- Improved class separation overall
- Motobike and traffic_light accuracy slightly better than with transfer learning
- Person detection more consistent
- Vehicle detection remains strong and comparable
- Confusion matrix shows cleaner boundaries between object classes and background

- Training time: 5h 30m (vs. 7h with transfer learning)
- Transfer learning required extra adaptation overhead

Conclusion and Further Work

Summary of Study

Evaluated multiple training strategies on a real-world object detection task using YOLOv8:

- Transfer Learning – Resized with label adjustment
- Transfer Learning – YOLO's Auto Resize & Padding
- Training from Scratch on Dataset 2

Key Findings

- Transfer learning improves generalization over no pretraining
- Manual resizing + corrected labels performs better than auto-padding for small objects
- Training from scratch gives best accuracy and fastest training time.
 - Especially for motobike and traffic_light classes

Conclusion

- YOLOv8s can be effectively trained from scratch on real-world data. Transfer learning helps, but is not always necessary when dataset size and quality are sufficient.

Future Directions

- Experiment with larger models: YOLOv8m, YOLOv8l



Thank you for
your attention

YOLOPilot